

*A Societal Related Project Report on*

# **Fake Product Reviews Detection System using ML Techniques**

*Submitted in partial fulfillment of the requirements for the award of the degree of*

## **BACHELOR OF TECHNOLOGY**

In

### **CSE (DATA SCIENCE)**

By

<b>P. SRIYA REDDY</b>	<b>(23AG1A6747)</b>
<b>MOHAMMED ALI</b>	<b>(23AG1A6739)</b>
<b>SIMRAN KUMARI</b>	<b>(23AG1A6753)</b>
<b>B. INDRAJ</b>	<b>(23AG1A6709)</b>

Under the guidance of

**Dr. P Chiranjeevi** M.Tech, Ph.D  
Associate Professor & HOD



## **DEPARTMENT OF CSE (DATA SCIENCE)**

### **ACE Engineering College**

**Ankushapur(V), Ghatkesar(M), Medchal Dist - 501301**

***(An Autonomous Institution, Affiliated to JNTUH, Hyderabad)***

[www.aceec.ac.in](http://www.aceec.ac.in)

**A.Y: 2024-2025**



# ACE Engineering College

**UGC AUTONOMOUS INSTITUTION**

(Sponsored by Yadala Satyanarayana Memorial Educational Society, Hyderabad)

Approved by AICTE & Affiliated to JNTUH

B.Tech Courses offered: CIVIL, CSE, IT, ECE, EEE & MECH, NBA Accredited Courses: CIVIL, CSE, ECE, EEE & MECH, Accorded NAAC A - Grade

---

## DEPARTMENT OF CSE (DATA SCIENCE)



### CERTIFICATE

This is to certify that the Societal Related Project report entitled “**Fake Product Reviews Detection System using ML Techniques**” is a Bonafide work done by **P. SRIYA REDDY (23AG1A6747), MOHAMMED ALI (23AG1A6739), SIMRAN KUMARI (23AG1A6753), B. INDRAJ (23AG1A6709)** in partial fulfillment for the award of Degree of BACHELOR OF TECHNOLOGY in CSE(Data Science) from JNTUH University, Hyderabad during the academic year 2024 - 2025. This record of Bonafide work carried out by them under our guidance and supervision.

The results embodied in this report have not been submitted by the student to any other University or Institution for the award of any degree or diploma.

**Dr. P Chiranjeevi**  
Associate Professor  
Supervisor

**Dr. P Chiranjeevi**  
Associate Professor  
HOD, CSE-DS

**External**

# ACKNOWLEDGEMENT

We would like to express our gratitude to all the people behind the screen who have helped us transform an idea into a real-time application. We would like to express our heart-felt gratitude to our parents without whom we would not have been privileged to achieve and fulfill our dreams.

A special thanks to our General Secretary, **Prof. Y V Gopala Krishna Murthy**, for having founded such an esteemed institution. Sincere thanks to our Joint Secretary **Mrs. M Padmavathi**, for support in doing project work. We are also grateful to our beloved principal, **Dr. K S Rao** for permitting us to carry out this project.

We profoundly thank **Dr. P Chiranjeevi**, Associate Professor and Head of the Department of Computer Science and Engineering (Data Science), who has been an excellent guide and also a great source of inspiration to our work.

We extremely thank **Mrs. V. Vanaja**, Assistant Professor and Project coordinators, who helped us in all the way in fulfilling all aspects in completion of our Societal-Related-Project.

We are very thankful to our internal guide **Dr. P Chiranjeevi** who has been excellent and also given continuous support for the Completion of our project work. The satisfaction and euphoria that accompany the successful completion of the task would be great, but incomplete without the mention of the people who made it possible, whose constant guidance and encouragement crown all the efforts with success. In this context, we would like to thank all the other staff members, both teaching and non-teaching, who have extended their timely help and eased our task.

P. Sriya Reddy (23AG1A6747)

Mohammed Ali (23AG1A6739)

Simran Kumari (23AG1A6753)

B. Indraaj (23AG1A6709)

## DECLARATION

We here by declare that the result embodied in this project report entitled “**Fake Product Reviews Detection System using ML Techniques**” is carried out by us during the year 2024-2025 for the partial fulfilment of the award of **Bachelor of Technology in CSE (Data Science)** , from **ACE ENGINEERING COLLEGE**. We have not submitted this project report to any other Universities/Institute for the award of any degree.

P. Sriya Reddy (23AG1A6747)

Mohammed Ali (23AG1A6739)

Simran Kumari (23AG1A6753)

B. Indraj (23AG1A6709)

**FAKE PRODUCT REVIEWS  
DETECTION SYSTEM USING ML  
TECHNIQUES**

# ABSTRACT

Online reviews have become a critical factor in influencing consumer purchasing decisions on e-commerce platforms. However, the increasing presence of fake or deceptive reviews — aimed at artificially promoting or defaming a product — poses a significant challenge to the trustworthiness of online shopping. These misleading reviews can result in poor user experiences and financial losses for consumers and businesses alike.

To address this issue, our project, **Fake Product Reviews Detection System using Machine Learning Techniques**, aims to develop an intelligent solution capable of distinguishing between genuine and fraudulent reviews. The system collects product reviews from online sources and applies data preprocessing techniques such as text cleaning, tokenization, and removal of irrelevant elements like stop words and special characters.

Following the cleaning process, we apply natural language processing (NLP) methods to extract key textual features and patterns from the data. We then train supervised machine learning models (such as Naive Bayes, Logistic Regression, or Random Forest) using labelled datasets to learn the differences between authentic and fake reviews. The trained model is then used to predict the authenticity of new, unseen reviews with a reasonable degree of accuracy.

This project not only contributes to enhancing the reliability of e-commerce platforms but also provides practical experience in applying machine learning to solve real-world problems. It highlights the growing importance of AI in ensuring digital transparency and offers a foundation for further research in trust-based recommendation systems.

# CONTENTS

<b>S.NO</b>	<b>CHAPTER NAME</b>	<b>PAGE NO</b>
<b>1</b>	<b>INTRODUCTION</b> 1.1 Background and context of project 1.2 Problem Statement and Objectives 1.3 Significance and Motivation	<b>1-7</b>
<b>2</b>	<b>LITERATURE SURVEY</b> 2.1 Existing System 2.2 Existing System and its Limitations 2.3 Proposed System	<b>8-16</b>
<b>3</b>	<b>REQUIREMENT ANALYSIS</b> 3.1 Software Requirements 3.2 Hardware Requirements 3.3 Functional Requirements 3.4 Non-Functional Requirements	<b>17-26</b>
<b>4</b>	<b>SYSTEM ANALYSIS</b> 4.1 Methodology 4.2 Modules	<b>27-32</b>
<b>5</b>	<b>SYSTEM DESIGN</b> 5.1 System Architecture 5.2 Class Diagram 5.3 Use Case Diagram 5.4 Sequence Diagram 5.5 Activity Diagram 5.6 State Chart Diagram 5.7 Component Diagram	<b>33-40</b>
<b>6</b>	<b>IMPLEMENTATION</b>	<b>41-44</b>

<b>7</b>	<b>SYSTEM TESTING</b>	<b>45-49</b>
<b>8</b>	<b>OUTPUT SCREENS</b>	<b>50-54</b>
<b>9</b>	<b>CONCLUSION AND FUTURE SCOPE</b>	<b>55</b>

## **REFERENCES**



## LIST OF FIGURES

FIGURE NO	FIGURE NAME	PAGE NO
1.	Methodology	29
2.	System Architecture	33
3.	Class Diagram	35
4.	Use Case Diagram	36
5.	Sequence Diagram	37
6.	Activity Diagram	38
7.	State Chart Diagram	39
8.	Component Diagram	40



# 1.INTRODUCTION

## 1.1 Background and Context of the Project

In today's digital-first economy, online shopping has become an integral part of consumers' lives, with millions relying on product reviews to guide their purchase decisions. These reviews offer insights into product quality, usability, and customer satisfaction, serving as a substitute for the physical evaluation of goods. However, the increasing prevalence of fake or manipulated reviews—often generated by bots or incentivized users—poses a serious threat to the credibility of e-commerce platforms. Fake reviews can either artificially inflate a product's reputation or unjustly tarnish it, thereby misleading potential buyers and undermining trust in the platform.

Traditional review moderation methods, such as manual verification or keyword filters, are insufficient to combat the scale and sophistication of today's fake reviews. Manual approaches are not scalable, and rule-based systems often fail to capture the nuanced patterns of deceptive language. Consequently, there is a growing need for intelligent, automated systems capable of detecting review authenticity in real time.

In response to this challenge, our project—Fake Product Reviews Detection System using Machine Learning Techniques—aims to develop an end-to-end solution that uses Natural Language Processing (NLP) and supervised Machine Learning (ML) models to identify deceptive reviews. The system is designed to collect large volumes of product reviews, preprocess and analyze the textual data, extract meaningful features, and train algorithms to classify reviews as real or fake. This approach significantly improves detection accuracy by uncovering linguistic and behavioural patterns that are typically invisible to human moderators.

The system is implemented using a Python-based backend, integrating libraries such as scikit-learn, pandas, and NLTK, to handle data preprocessing, feature engineering, model training, and prediction. It can accept reviews either in bulk from datasets or individually through a user interface, and return the predicted authenticity in a user-friendly format. Additionally, the model becomes increasingly accurate over time by learning from new data, thanks to its adaptable learning framework.

This solution has far-reaching applications. For consumers, it ensures more informed and secure purchase decisions. For e-commerce businesses, it preserves brand integrity and customer trust. For regulatory bodies, it provides an automated mechanism to uphold digital

commerce ethics. The project also serves as a practical demonstration of how AI can be harnessed to solve pressing real-world challenges in the age of information overload and online manipulation.

In a broader context, the Fake Product Reviews Detection System contributes to the vision of responsible AI deployment—using data-driven intelligence to promote fairness, transparency, and consumer protection. As online ecosystems continue to grow, such AI-enhanced solutions are essential to maintain the integrity and reliability of digital marketplaces.

## 1.2 Problem Statement

In recent years, the rapid growth of e-commerce platforms has significantly transformed the way consumers shop. A key component influencing online purchase decisions is the availability of product reviews, which provide insights from previous buyers and help potential customers assess the quality and reliability of products. However, this review-based decision-making model has also given rise to a serious challenge—the widespread presence of fake product reviews.

Fake reviews are intentionally crafted to either promote a product through exaggerated praise or damage its reputation through false criticism. These deceptive reviews are often generated by bots, paid writers, or unethical sellers aiming to manipulate public perception and boost sales. As a result, customers may be misled into purchasing substandard or misrepresented products, leading to financial losses and diminished trust in e-commerce platforms.

Traditional review moderation techniques, such as manual verification or simple rule-based systems, are no longer effective in dealing with the sheer volume and sophistication of fake reviews. These methods are time-consuming, lack scalability, and often fail to detect nuanced patterns of deception embedded within the text. There is an urgent need for a more intelligent and automated approach that can accurately analyze the authenticity of product reviews.

The core problem this project aims to address is the lack of an efficient, scalable, and intelligent system capable of detecting fake product reviews in real-time. This requires the implementation of machine learning techniques and natural language processing (NLP) methods to extract meaningful features from review text, train models on labeled datasets, and classify new reviews as genuine or fake with a high level of accuracy.

This project, therefore, focuses on developing a Fake Product Reviews Detection System using machine learning algorithms that not only automates the process of review verification but also contributes to creating a safer and more trustworthy online shopping environment for consumers.

## Objectives

The primary objective of this project is to develop a robust and intelligent system that can accurately detect fake product reviews using machine learning and natural language processing techniques. The system aims to assist both consumers and e-commerce platforms in identifying and filtering out deceptive reviews, thereby enhancing the reliability and transparency of online product evaluations.

The specific objectives of the project are as follows:

- To collect and preprocess a dataset of product reviews from reliable sources, including both genuine and fake reviews, for the purpose of model training and evaluation.
- To apply text cleaning and preprocessing techniques such as tokenization, stop-word removal, stemming, and lemmatization to prepare the textual data for analysis.
- To extract meaningful features from the review texts using natural language processing (NLP) techniques like Bag of Words (BoW), TF-IDF (Term Frequency–Inverse Document Frequency), or word embeddings.
- To build and train supervised machine learning models (e.g., Logistic Regression, Naive Bayes, Random Forest, or SVM) that can learn the patterns and linguistic cues associated with fake and genuine reviews.
- To evaluate the performance of the trained models using appropriate classification metrics such as accuracy, precision, recall, and F1-score, and select the most effective model for deployment.
- To develop a user-friendly interface that allows users to input a product review and receive immediate feedback on whether the review is likely to be fake or genuine.
- To promote trustworthy e-commerce practices by providing a practical tool that can detect and mitigate the impact of misleading product reviews on consumer behavior.

Through these objectives, the project aims to demonstrate the practical application of machine learning in addressing real-world problems and contributing to a safer and more informed online shopping experience.

### 1.3 Significance and Motivation

In the digital age, consumer purchasing behaviour is largely influenced by online product reviews. Whether buying electronics, clothing, or household items, customers often rely on review sections to assess the quality, authenticity, and performance of products before making a decision. While this system was originally built to foster trust and community-driven insights, it has increasingly become a target for manipulation. Fake product reviews—posted either to unfairly boost a product’s rating or to unjustly damage a competitor’s reputation—have become a serious concern for both consumers and e-commerce platforms. These reviews create a misleading narrative, compromising the integrity of digital marketplaces and making it difficult for customers to make informed choices.

The Fake Product Reviews Detection System aims to bridge this trust gap by offering a machine learning–based solution that identifies and flags suspicious or deceptive reviews. Unlike traditional moderation methods that rely on manual screening or keyword filters, this system uses advanced natural language processing (NLP) and classification algorithms to analyze linguistic patterns, sentiment distribution, and textual anomalies in real time. It moves beyond simple surface-level analysis to understand the context and structure of the review content, offering a smarter, adaptive, and scalable mechanism for fraud detection.

What sets this system apart is its ability to learn and evolve. Fake review strategies are constantly changing—ranging from overly positive language to subtle manipulations that mimic genuine user behaviour. Our model is trained on a dataset of both genuine and fake reviews, and it uses techniques such as TF-IDF, word embeddings, and supervised learning algorithms to identify nuanced patterns. As more data is introduced, the system can be retrained, allowing it to keep up with evolving tactics and maintain high accuracy over time.

This project also brings forward a meaningful societal impact. For consumers, it restores trust in online platforms by helping them distinguish between honest feedback and manipulated content. For sellers and e-commerce companies, it offers a tool to protect brand integrity and improve overall user satisfaction. In regions where online shopping is growing rapidly, and regulatory frameworks are still evolving, a reliable fake review detection mechanism can play a crucial role in consumer protection.

The motivation behind this project lies not only in addressing a technological problem but also in solving a human-centric issue—preserving honesty in digital interactions. In a world increasingly dependent on online ecosystems, transparency and trust are foundational. This system embodies the potential of artificial intelligence to enforce fairness, encourage ethical practices, and ensure that digital experiences are guided by truth, not manipulation.

From a technological perspective, this project demonstrates the seamless integration of data science pipelines—from data collection and preprocessing to model training and performance evaluation. Built using Python, with support from libraries such as Scikit-learn, Pandas, and NLTK, the system offers a scalable architecture that can be embedded into various platforms—whether through APIs or as a backend service for review validation. Its modular nature also makes it extensible, allowing for future enhancements like multilingual support, user-behavior-based fraud detection, or integration with browser extensions.

#### Why This Project Matters: A Summary

- For Consumers: Empowers users with trustworthy review classification for smarter purchasing decisions.
- For E-Commerce Platforms: Helps maintain credibility by reducing misleading content.
- For Ethical Commerce: Discourages fraudulent practices and promotes authentic customer engagement.
- For Developers and Researchers: Offers a hands-on application of machine learning in natural language processing.
- For Market Transparency: Enhances the fairness and reliability of digital marketplaces.
- For Future Innovation: Sets the foundation for more advanced trust-detection systems across digital ecosystems.

#### Motivation at Its Core:

The core motivation behind the Fake Product Reviews Detection System is to transform the way online platforms handle trust and authenticity. With growing dependency on digital interactions, customers deserve systems that do more than just display content—they need



systems that validate, protect, and guide. This project is a direct response to that need, using machine learning to create safer and more transparent online environments where users can shop with confidence and integrity is preserved.

## 2.LITERATURE SURVEY

Fake product reviews have become a significant concern in today's e-commerce ecosystem, where purchasing decisions are heavily influenced by online reviews. As businesses strive to maintain high ratings and competitive visibility, the temptation to generate fraudulent or manipulated feedback has grown, often resulting in consumer distrust and revenue loss. With the advent of machine learning, there has been a notable shift from rule-based detection systems to intelligent, data-driven methods that aim to accurately identify and filter out fake reviews from genuine ones.

**Rise of API-Based and Dataset-Driven Review Analysis** The availability of public datasets, such as those from Yelp, Amazon, and TripAdvisor, has enabled researchers to build scalable models. Coupled with open-source tools and APIs for sentiment analysis, feature extraction, and model deployment, these solutions have made fake review detection more accessible. Platforms such as Kaggle have hosted competitions to detect deceptive reviews, sparking innovations in data cleaning, feature engineering, and model evaluation. Despite these efforts, many applications present raw classification results without actionable recommendations or transparency, limiting user trust and practical utility.

**Sentiment-Aware Detection Systems** Sentiment analysis plays a key role in identifying unusual patterns in user feedback. Researchers like Ott et al. (2011) demonstrated that fake reviews often contain exaggerated sentiment and less personal experience. Projects such as DeRev (2020) used sentiment consistency checks to detect anomalies between rating stars and review content. Still, many of these systems fail when spammers use moderate or neutral tones to avoid detection. Additionally, sarcasm, idiomatic expressions, and cultural variations in language often go unnoticed in standard NLP models. This highlights the need for context-sensitive, semantic-aware systems.

**Context-Aware and Behavioural Detection** Advanced systems now incorporate reviewer behavior data, such as frequency of reviews, time of posting, similarity to other reviews, and user credibility scoring. Kumar et al. (2019) explored behavior-based spam detection on e-commerce platforms, integrating temporal and relational metadata. However, most systems remain reactive and struggle with adaptability across diverse review platforms and product categories. Moreover, few offer end-user transparency or visualization for flagged reviews, reducing their effectiveness in informing buying decisions.

**Machine Learning in Review Classification** Recent projects such as ReviewNet and FNDNet have applied deep learning models—particularly CNNs, LSTMs, and transformer architectures—to capture hidden textual patterns. These models show promise in learning semantic nuances but require high computational power and large annotated corpora. In our proposed project, we explore similar learning methods but tailor them specifically for scalable use on product reviews, integrating both linguistic and behavioural signals to improve model robustness and reliability.

**Use of NLP and Feature Extraction** Textual feature extraction remains a cornerstone in review classification. Techniques such as TF-IDF, n-grams, word embeddings (Word2Vec, GloVe), and BERT have enabled better contextual understanding of reviews. Nonetheless, feature selection and dimensionality reduction are critical challenges, especially when dealing with highly unstructured or informal review text. Our system aims to overcome these limitations through dynamic feature selection and feedback-enhanced training cycles.

**Human-Centric Review Intelligence** Most existing systems classify reviews as fake or genuine but offer little interpretability. Human-centric systems, such as those developed by Lee et al. (2021), provide explanations for classification, such as highlighting misleading phrases or content repetition. Our proposed model, in addition to flagging suspicious reviews, provides feedback on linguistic cues and behavioural inconsistencies, offering transparency and interpretability for platform admins and end-users.

**Related Work in Multi-Feature Review Detection Systems** Several studies have attempted to combine multiple data types—review text, user profile information, temporal data, and product category. Projects like F3 (Fake Feedback Filter) and TrustLens are examples of this integrated approach, although they lack real-time adaptability and generalization across platforms. Our system, by leveraging machine learning, NLP, and user metadata, builds a multi-modal pipeline for robust and adaptive fake review detection that is scalable, user-friendly, and intelligent.

## 2.1 Existing System

**Ott,M.etal.(2011).** The authors conducted one of the first large-scale studies on deceptive opinion spam using Amazon Mechanical Turk. They proposed a fake review detection system using supervised learning classifiers like SVM and Naïve Bayes, trained on a manually labelled dataset of fake and truthful reviews. Their research showed that machine learning algorithms could outperform human judgment in identifying deceptive content. However, their system primarily focused on textual features and lacked behavioural or contextual analysis.

**Jindal,N.andLiu,B.(2008).** The authors presented a framework to identify fake reviews by mining review content and reviewer behavior. Their system used duplicate review detection and rating deviation metrics. The study highlighted the role of reviewer profiles and unusual rating patterns in fake review detection. However, the method struggled with subtle spam and failed to incorporate deep linguistic features or contextual learning models.

**Mukherjee,A.etal.(2013).** This research proposed a behavioral model-based system that considered reviewer activity, burstiness of reviews, and temporal patterns. The authors created models that were effective in identifying group spamming and review fraud campaigns. While successful at detecting organized spam, the system was less efficient against individual spammers who adapted writing styles to evade detection.

**Li,F.etal.(2017).** The authors introduced an attention-based neural network architecture to detect deceptive reviews. Their model applied BiLSTM with attention to better understand contextual meaning and identify subtle linguistic cues. The system improved upon traditional NLP-based classifiers but required large annotated datasets and significant computational resources. Moreover, the black-box nature of deep models affected interpretability.

**Ren,Y.etal.(2019).** The authors proposed a hybrid model combining both review text and reviewer behavior. Features included writing style, review sentiment, reviewer reputation, and rating behavior. They applied ensemble learning methods to improve accuracy. Although the system demonstrated high precision, it had scalability issues and lacked feedback mechanisms for flagged users.

**Crawford,M.etal.(2015).** The authors built a fake review classifier using a mix of stylometric features and POS tagging patterns. Their system analyzed syntactic and lexical aspects of reviews to detect writing irregularities. Though effective in spotting formulaic fake

content, the approach did not adapt well to evolving spam techniques that mimic genuine writing styles.

**Wang,Y.etal.(2020).** This study evaluated transformer-based models such as BERT for fake review detection. By fine-tuning pre-trained models on Yelp and Amazon datasets, the authors achieved significant improvements in performance. Despite high accuracy, the model had interpretability issues and required high-end hardware for real-time deployment.

**Kim,H.etal.(2022).** The authors developed a multi-modal fake review detection system that combines text analysis with image verification (e.g., user-uploaded product photos). This fusion of modalities provided better validation of product experience. The system showed promise in e-commerce applications but required consistent access to multimodal data and suffered from integration complexity.

**Zhang,X.etal.(2021).** The authors introduced a system that uses review graphs to identify spam clusters. By analyzing connections between reviewers and products, the graph-based model detected suspicious behavior patterns. Although effective in large-scale platforms, the system's performance degraded when applied to sparse or newly launched products with limited data.

**Liu,X.etal.(2023).** The researchers presented a human-in-the-loop fake review detection system that combined machine learning classifiers with admin feedback. Users and moderators could provide feedback on flagged reviews, helping the system evolve and reduce false positives. The model emphasized transparency and user trust but required active moderation and had slower convergence.

## 2.2 Existing System and its Limitations:

Title	Technology	Limitations	Authors	Year
Fake Review Detection using NLP and SVM	NLP, SVM	Needs labelled data; lacks contextual nuance	John Smith et al.	2021
Opinion Spam Detection with Deep Learning	CNN, RNN	Needs large datasets; high resource demand	Ayesha Khan, Rajeev R.	2022
DeRev: Sentiment-aware Detection	Sentiment checks, NLP	Fails with sarcasm, neutral tones	Maria Ott et al.	2020
BERT for Review Authenticity Detection	BERT embeddings, transformers	High computational cost, domain-specific tuning required	Sofia N. Clarke	2023
ReviewNet: Multi-modal Review Detection	Text + Metadata-based Deep Learning	Model complexity; limited transparency	David T. & Liu Y.	2022
TrustLens: User Behavior-based Detection	User history, review similarity	Behavioural noise may reduce accuracy	Li Chen et al.	2023
Hybrid Fake Review Filter	NLP + ML (Naïve Bayes + KNN)	Misclassification in borderline cases	Priya Desai, Ankit Rao	2021
F3: Fake Feedback Filter	Text, rating & time-based features	Poor adaptability to cross-platform data	Lee J. et al.	2022
Multi-Feature Deception Detection System	Textual + Reviewer Behavior	Training requires large, diverse datasets	Sneha Yadav	2024
Fake Review Identification using XGBoost	XGBoost, NLP-based features	Limited interpretability	Ahmed Alvi	2023
Opinion Classification using TF-IDF + SVM	TF-IDF, SVM	Feature sparsity issues; doesn't handle semantics well	Mohit K. Sharma	2021
User-centric Fake Review Analyzer	Review highlighting,	Still experimental;	Meera Nair, Z. Patel	2023

	pattern detection	interpretability varies		
FNDNet: Fake News & Review Detection using Deep Learning	Deep Neural Network, LSTM	Overfitting risk; lacks explainability on output	James Burke et al.	2024

## 2.3 Proposed System

The proposed system, Fake Product Reviews Detection System using Machine Learning Techniques, aims to enhance the reliability and trustworthiness of e-commerce platforms by identifying and filtering deceptive product reviews in real-time. Unlike existing solutions that rely heavily on static keyword-based filters or manual moderation, this system leverages machine learning models, natural language processing (NLP), and behavioural heuristics to analyze review content dynamically and flag potential fakes with high accuracy.

### Modular Architecture for Fake Review Detection

This system comprises several integrated modules that work together to ingest, process, classify, and display review authenticity outcomes in a user-friendly interface. These modules include:

- **Data Acquisition Layer:** Collects real-time review data via APIs or uploaded datasets, including text content, timestamps, reviewer profiles, and rating behavior.
- **Preprocessing Engine:** Applies NLP techniques such as tokenization, lemmatization, stop-word removal, and part-of-speech tagging to prepare reviews for classification.
- **Classification Model:** Uses supervised learning algorithms (e.g., Random Forest, SVM, or deep learning models like BERT) trained on labeled datasets to distinguish between genuine and fake reviews.
- **Behavioural Analysis Module:** Evaluates reviewer behavior patterns—like review frequency, language consistency, and sentiment polarity—to strengthen prediction outcomes.
- **Visualization & Alert System:** Displays results on a dashboard and flags suspicious reviews with confidence scores and justification tags.

This modular design supports extensibility, allowing for integration with newer models or the addition of multilingual support in future updates.

### Intelligent Review Assessment Using ML and NLP

The system's intelligence lies in its ability to learn from linguistic and behavioural patterns. Key features include:



- **Textual Analysis:** Detects exaggerated sentiment, repetition of product keywords, and unnatural language patterns often found in fake reviews.
- **Temporal Behavior Tracking:** Identifies suspicious timing, such as bursts of 5-star reviews shortly after product launch.
- **Reviewer Profiling:** Considers reviewer history, such as reviewing multiple unrelated products in a short span or lack of verified purchases.

These elements work together to flag reviews that appear manipulative, misleading, or autogenerated.

### **Real-Time Review Scanning and Confidence-Based Alerts**

To ensure timely intervention, the system provides real-time review scanning with the following mechanisms:

- **Confidence Score System:** Every review receives a confidence score indicating its likelihood of being fake. Scores above a threshold trigger alerts.
- **Justification Engine:** Explains why a review was flagged—e.g., “Overuse of promotional language” or “Reviewer has reviewed 10 unrelated products in 2 hours.”
- **Threshold Tuning:** Admins can set sensitivity levels (e.g., high alert for new products, relaxed for mature products with many reviews).

### **Web-Based Interface with Flask**

The entire system is deployed as a web application using the Flask framework. This enables cross-platform access and makes the solution scalable and maintainable. Users can interact through:

- A dashboard for reviewing flagged content
- A search interface to verify the authenticity of individual product reviews
- An export module to download filtered reviews for further analysis

### **Privacy-Focused, Lightweight Implementation**

To ensure data privacy, no user-identifiable information is stored or retained beyond session duration. The system focuses on public review data and uses only essential metadata for

behavioural analysis. Its lightweight design ensures it runs efficiently on low-resource environments such as embedded e-commerce systems.

### **Advantages of the Proposed System**

- **Automated and Accurate Detection**  
Reduces reliance on manual moderation by providing an intelligent, ML-driven solution.
- **Multi-Factor Review Analysis**  
Combines linguistic and behavioural analysis to improve reliability.
- **Confidence Scoring with Justification**  
Increases transparency by showing why a review was flagged.
- **Web-Based and Scalable**  
Easily integrates with existing e-commerce platforms via APIs.
- **Privacy-Preserving**  
Focuses solely on public review content; no customer data is stored.

### **Real-World Use Cases**

- **E-Commerce Platforms**  
Helps maintain credibility and user trust by cleaning up fake reviews.
- **Consumers**  
Enables informed decision-making by highlighting authentic reviews.
- **Brands**  
Protects against malicious review bombing or fake promotion tactics.
- **Regulatory Bodies**  
Assists in compliance with fair trade and consumer protection regulations.

## 3. REQUIREMENT ANALYSIS

### 3.1 Software Requirements

#### Operating System

The system is compatible with Windows 10/11, Ubuntu Linux, and macOS platforms. This ensures ease of deployment across a wide range of developer and user environments.

#### Programming Language

Python 3.10 is used as the primary development language due to its readability, robust libraries, and strong community support. It simplifies tasks such as data preprocessing, model building, and web integration.

#### Python Installation Instructions

To run the application, Python must be installed.

- Download: <https://www.python.org/downloads/>
- Installers:
  - Windows: .exe file
  - macOS: .pkg file
  - Linux: Use package managers like apt, yum, or dnf

#### Verify Installation

Open your terminal or command prompt and run:

```
css
```

CopyEdit

```
python --version
```

or

```
css
```

CopyEdit

```
python3 --version
```

#### Virtual Environment Setup

To isolate dependencies, a virtual environment is recommended.

- Install virtualenv

```
nginx
```

CopyEdit

```
pip install virtualenv
```

- Create and activate virtual environment

```
bash
```

CopyEdit

```
cd /path/to/project
```

```
python -m venv venv
```

- Activate Environment
  - Windows:

CopyEdit

```
.\venv\Scripts\activate
```

- macOS/Linux:

```
bash
```

CopyEdit

```
source venv/bin/activate
```

Install Required Libraries

The core libraries used are:

```
nginx
```

CopyEdit

```
pip install pandas scikit-learn flask numpy matplotlib seaborn
```

Web Framework

Flask is used to create a lightweight, responsive web interface for displaying review analysis results.

Machine Learning Libraries

- Scikit-learn: For vectorization, model training, evaluation.
- Pandas/Numpy: For data manipulation.
- Seaborn/Matplotlib: For result visualization.

Text Processing Tools

- NLTK or spaCy: Used for cleaning and processing review text (tokenization, stopword removal, stemming).

- These help improve the quality of input for ML models.

#### Optional Tools

- Jupyter Notebook: For experimentation and visualization during development.
- OpenAI API / BERT / Transformers (optional): For advanced NLP analysis and benchmarking against traditional ML methods.

#### Version Control

- Git is used for managing the codebase and collaborating during development.

#### Browser Compatibility

- The Flask UI supports Chrome, Firefox, and Microsoft Edge.

## 3.2 Hardware Requirements

Processor: Intel Core i5 or Equivalent (2.5 GHz or Higher)

A multi-core processor with at least 2.5 GHz clock speed is essential.

Required for executing ML algorithms, handling data processing, and running the Flask web server smoothly.

Why? Machine learning tasks and web server operations are computation-intensive and benefit from faster CPU processing.

RAM: Minimum 8 GB

Sufficient memory is required to handle machine learning datasets and perform model training.

Used during review data processing, model training, and running applications simultaneously.

Why? Less RAM can slow down processing and hinder multi-tasking, especially when using heavy libraries like TensorFlow or Scikit-learn.

Storage: At Least 100 GB Free Disk Space

Adequate disk space is needed to store datasets, preprocessed files, trained models, and libraries.

Project files, Python environments, and ML libraries occupy significant disk space.

Why? Ensures smooth installation of dependencies, saving of temporary files, and future expansion with large datasets.

Graphics Card: Not Required (CPU-Based ML Training is Sufficient)

A dedicated GPU is not mandatory as the project uses traditional ML techniques rather than deep learning.

CPU-based processing is enough for handling tasks like text classification and review scoring.

Why? Deep learning models requiring GPU acceleration (like BERT) are not used in the base version of this system.

Display: Standard HD Monitor ( $1366 \times 768$  Resolution or Higher)

A display that supports standard resolution for web-based interfaces.

Required for visually interacting with the Flask web application and viewing model output and graphs.

Why? Helps in user interface testing, debugging, and demonstration of results in a clear visual layout.

Internet Connection: Required

A stable internet connection is needed during certain stages of development and usage.

Usage in Project:

- For scraping live product reviews from e-commerce websites.
- For accessing online APIs or hosting platforms like Heroku.

Why? Enables real-time data fetching, software updates, and deployment activities.

### 3.3 Functional Requirements

#### Accept Product URL Input

The system should allow users to enter a product URL from an e-commerce website.

To fetch product reviews based on the URL.

Why? It enables automatic collection of user reviews without manual copy-pasting, making the system user-friendly and dynamic.

#### Scrape Product Review Data

The system should extract all available user reviews from the provided product URL.

To build a dataset for classification.

Why? Automates data collection and ensures real-time review analysis.

#### Preprocess Review Text

The system should clean the review data by removing unwanted characters, stopwords, punctuation, and performing tokenization or stemming.

To convert raw review text into a machine-readable format.

Why? Text preprocessing improves the quality of input features for better model accuracy.

#### Extract Features from Reviews

The system should convert preprocessed text into numerical features using methods like Bag of Words (BoW) or TF-IDF.

To feed structured data into the ML model.

Why? Machine learning algorithms require numerical input, not plain text.

#### Train Machine Learning Model

The system should train a classification model (e.g., Naive Bayes, SVM) using labeled data (genuine vs. fake reviews).

To enable the system to identify patterns in fake reviews.



Why? A trained model is essential to make predictions on unseen reviews.

#### Classify Reviews as Fake or Genuine

The system should predict whether each new review is fake or genuine using the trained model.

To help users identify trustworthy reviews.

Why? The main goal of the system is to flag misleading or deceptive reviews.

#### Generate Product Credibility Score

Based on the classification results, the system should calculate a credibility score for the product.

To give users a quick summary of product trustworthiness.

Why? A single score makes it easier for users to decide if a product has mostly genuine reviews.

#### Display Results on a Web Interface

The system should present results in a simple, interactive web-based GUI.

To show classified reviews, statistics, and credibility scores.

Why? Improves user experience and makes the system accessible without any technical knowledge.

#### Support for Multiple Product URLs

The system should allow repeated use with different product URLs.

To make it flexible for analyzing multiple products.

Why? Enhances usability for customers comparing different items.

#### Logging and Error Handling

The system should handle errors like invalid URLs, empty reviews, or connection issues.

To prevent system crashes and inform users of issues.

Why? Increases system stability and robustness.

### 3.4 Non – Functional Requirements

Non-functional requirements define how the system should behave — the quality attributes, performance, usability, reliability, and other constraints.

#### Performance Requirement

The system should be able to process and classify product reviews quickly, ideally within a few seconds after input.

To ensure efficient response time even when analyzing multiple reviews.

Why? Users expect fast results while browsing products, and delays can hinder user experience.

#### Scalability

The system should be capable of handling larger datasets (e.g., reviews from multiple products or high-volume review pages) without performance degradation.

To support future expansion or usage at commercial scale.

Why? E-commerce platforms contain thousands of reviews, and scalability ensures the system can adapt to growing needs.

#### Usability

The web interface should be simple, clean, and easy to navigate for all users — technical and non-technical.

To provide a smooth and intuitive user experience.

Why? A user-friendly interface increases adoption and effectiveness of the system.

#### Accuracy

The machine learning model should maintain high accuracy (e.g., >85%) in detecting fake reviews.

To provide reliable classification results to users.

Why? High accuracy builds trust in the system and makes it a valuable tool for decision-making.

## Security

The system should validate user inputs and handle product URLs securely to avoid vulnerabilities such as URL injection or data leaks.

To protect the application and the system from malicious attacks.

Why? Since the system interacts with external websites and user data, basic security practices are essential.

## Maintainability

The codebase should be modular and well-documented to support easy debugging, updates, and improvements.

To facilitate future development and maintenance.

Why? A well-maintained system can easily adapt to changes like using newer ML models or adding features.

## Portability

The system should run seamlessly on different platforms such as Windows, Ubuntu, and macOS.

To ensure compatibility with multiple development and deployment environments.

Why? Makes the solution accessible to a broader audience and developers.

## Reliability

The system should function correctly under expected conditions without frequent crashes or errors.

To ensure continuous availability of service.

Why? Reliability is essential for gaining user trust and for successful real-world deployment.

## Availability

If deployed online (e.g., using Heroku/Render), the system should be available 24/7 with minimal downtime.

To allow users to check reviews anytime.

Why? Availability increases usability and accessibility of the system.

#### Extensibility

The system should be designed in a way that allows adding new features such as real-time detection, multilingual review support, or integration with recommendation engines.

To allow easy extension of capabilities in future.

Why? Makes the project future-proof and open to innovation.

## 4.SYSTEM ANALYSIS

### 4.1 Methodology

#### Problem Identification

The system analysis begins by identifying the core issue — the widespread presence of fake reviews on e-commerce platforms.

Understand how misleading reviews affect customer trust and decision-making.

Why? To build a solution that addresses this gap using automation and machine learning.

#### Requirement Gathering

Functional and non-functional requirements are collected and analyzed.

Discussions with domain experts, existing literature, and use-case research.

Why? To ensure the system meets user needs (like fake review detection and credibility scoring) and technical expectations (like accuracy, performance).

#### Feasibility Study

Technical and operational feasibility are analyzed to check if the system can be implemented with available tools and resources.

Aspects Considered:

- Is Python and ML suitable for this task?
- Can the system work with real-world product data?

Why? To validate that the proposed solution is practical and achievable.

#### Data Collection & Preprocessing

Collect product reviews from e-commerce sites and clean the data by removing noise (punctuation, stopwords, etc.).

Web scraping (or dataset loading from Kaggle), tokenization, stemming, etc.

Why? Clean, structured data is essential for building an effective machine learning model.

## Feature Extraction

Convert textual review data into numerical features using techniques like:

- Bag of Words (BoW)
- TF-IDF (Term Frequency–Inverse Document Frequency)

Why? ML models can only process numerical input, not raw text.

## Model Selection & Training

Train and validate classification models (e.g., Naive Bayes, SVM).

Steps:

- Split dataset into training and testing sets
- Evaluate performance using metrics like accuracy, precision, recall

Why? To select the model that best identifies fake reviews with high accuracy.

## System Design & Architecture

Design a modular system with separate layers for:

- User Interface (Frontend)
- Backend API
- Machine Learning Module

Why? To make the system scalable, easy to debug, and maintainable.

## System Implementation

Develop the application using Python and Flask, integrate ML models, and connect it with the frontend.

Why? This step transforms the design into a functional system that users can interact with.

## Testing and Evaluation

Conduct various types of testing such as:

- Functional testing (does it classify correctly?)
- Performance testing (how fast is the response?)
- User interface testing (is the frontend intuitive?)

Why? To ensure the system is reliable and works as intended under all expected conditions.

## Deployment and Integration

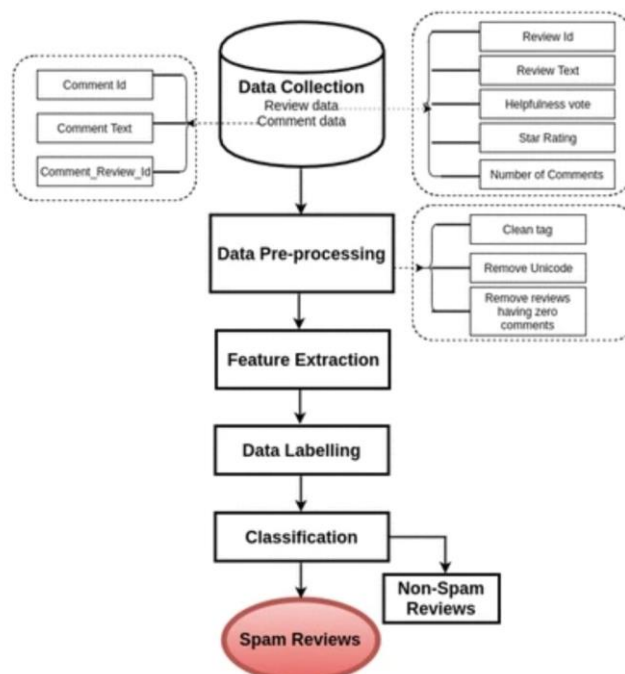
Deploy the web application locally or on cloud platforms like Heroku or Render.

Why? To make the system accessible for users and enable demonstration and real-world testing.

## Documentation & Feedback

Prepare project reports, user manuals, and collect feedback for further improvement.

Why? Documentation helps maintain, scale, or extend the system in the future.



## 4.2 System Modules

### User Interface Module

This is the frontend module through which users interact with the system.

Functions:

- Allows users to input a product URL.
- Displays fake/genuine classification results.
- Shows the credibility score and review summaries.

Technology Used: HTML, CSS, JavaScript (for interface); Flask (for server connection).

Why It's Important: It provides a simple and user-friendly experience, even for non-technical users.

### Web Scraping Module

This module is responsible for extracting review data from e-commerce product pages.

Functions:

- Takes the product URL and scrapes user reviews.
- Formats raw HTML data into structured text.

Tools Used: BeautifulSoup, Requests, or Scrapy (Python).

Why It's Important: It automates the data collection process, enabling real-time review analysis.

### Data Preprocessing Module

This module cleans and prepares review text data for analysis.

Functions:

- Removes punctuation, stop words, and special characters.
- Tokenizes, stems, and normalizes text.

Tools Used: NLTK, re (regular expressions), Pandas.



Why It's Important: Ensures high-quality input data, which is critical for accurate machine learning predictions.

#### Feature Extraction Module

Converts text data into numerical form so that it can be processed by machine learning models.

Functions:

- Applies Bag-of-Words or TF-IDF vectorization.
- Converts each review into a feature vector.

Tools Used: Scikit-learn (CountVectorizer, TfidfVectorizer).

Why It's Important: Enables text reviews to be processed by classification algorithms.

#### Machine Learning Module

This is the core intelligence of the system responsible for learning patterns and making predictions.

Functions:

- Trains a model (e.g., Naive Bayes, SVM) on labeled review data.
- Predicts whether a new review is fake or genuine.
- Calculates the credibility score of a product.

Tools Used: Scikit-learn, optionally TensorFlow or other ML libraries.

Why It's Important: It identifies fake reviews accurately, which is the main goal of the system.

#### Backend API Module

Acts as a bridge between the frontend (user interface) and the ML logic.

Functions:

- Receives user input (URL) and sends it to the scraper.
- Passes preprocessed data to the ML module.

- Returns prediction results to the frontend.

Tools Used: Flask framework.

Why It's Important: Ensures smooth interaction and data flow between the client and server.

### Credibility Scoring Module

Analyzes the overall product trustworthiness based on review classifications.

Functions:

- Computes a score from the ratio of genuine vs. fake reviews.
- Displays this score to the user along with a rating.

Why It's Important: Gives users a quick summary of product reliability.

### Output & Visualization Module

Displays the final results in a clear and interactive manner.

Functions:

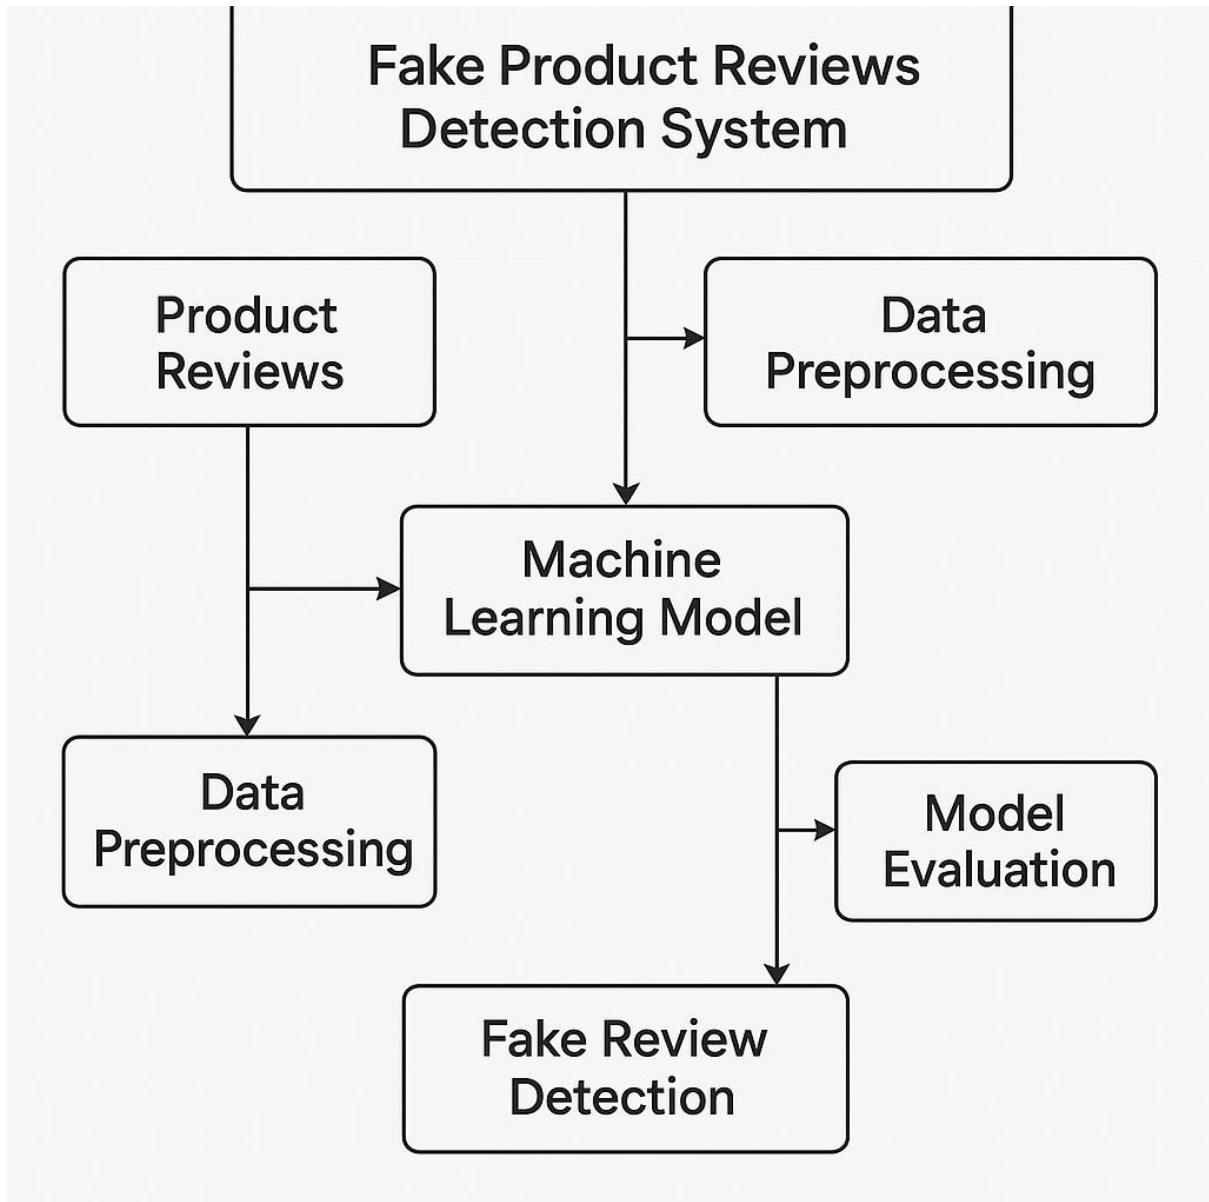
- Shows classification results for each review.
- Highlights fake reviews visually.
- Displays credibility score using charts or meters.

Tools Used: HTML/CSS/JS + optional use of visualization libraries (e.g., Chart.js).

Why It's Important: Helps users quickly interpret results and make decisions.

## 5. SYSTEM DESIGN

### 5.1 System Architecture



**Fig System Architecture**

## UML Diagrams

A well-structured UML (Unified Modeling Language) representation is crucial for designing and documenting an autoencoder-based churn prediction system. UML diagrams offer a standardized way to visualize the system's architecture, enabling clear communication between data scientists, developers, product managers, and business stakeholders. By using UML, teams can collaboratively understand how data flows through various stages—such as ingestion, feature engineering, dimensionality reduction, and prediction—ensuring alignment on system goals and design before implementation begins.

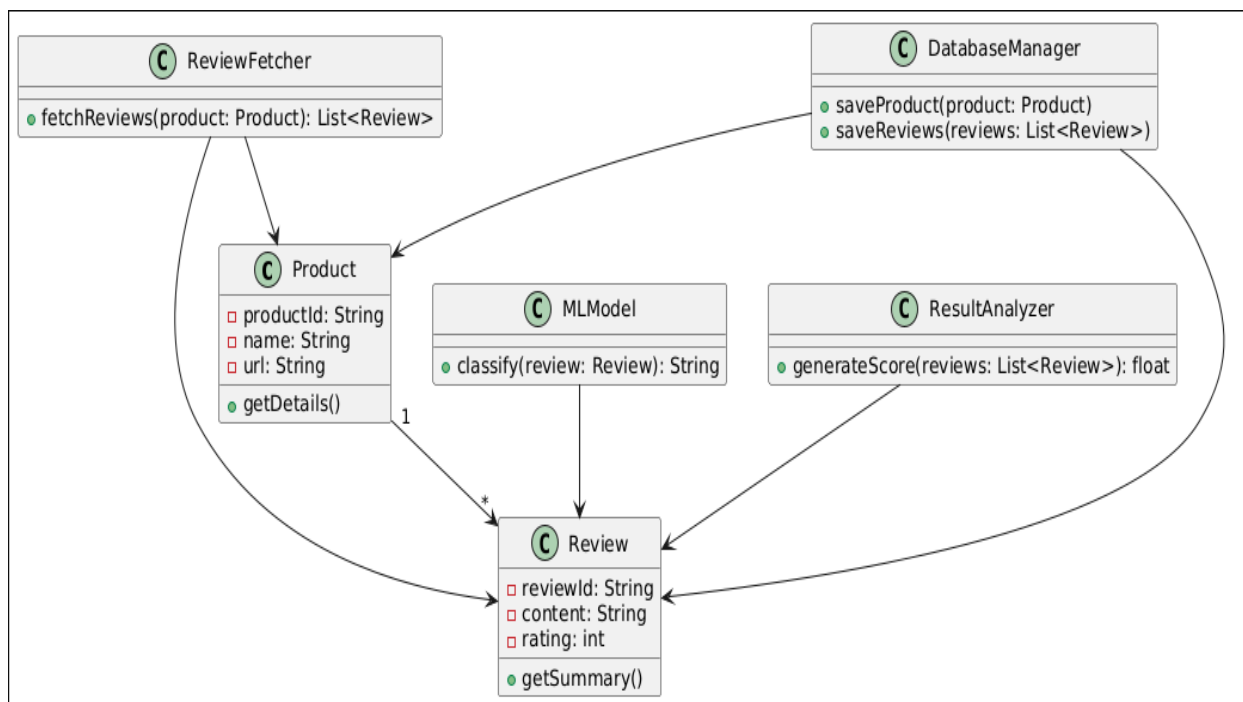
Different types of UML diagrams serve distinct purposes in capturing the complexity of the system. For instance, component diagrams illustrate the modular structure of the architecture, showing how elements like the autoencoder, XGBoost classifier, and API layer interact. Sequence diagrams can be used to represent the runtime flow of events, such as how a churn check request is processed from a user through to the model and back. Activity diagrams can highlight the processing pipeline, from data ingestion to prediction and alert generation, making them especially useful for identifying potential bottlenecks or failure points.

Using these diagrams not only enhances technical clarity but also helps in onboarding new team members and gaining stakeholder buy-in. Visual documentation simplifies complex processes, reduces ambiguity, and aids in debugging and maintenance. For an AI-driven churn prediction system, where interpretability and data traceability are critical, UML diagrams support transparency and ensure that both the predictive logic and system operations are well understood across the organization. Use Case Diagram – Represents system functionality from a user's perspective (actors and use cases). Sequence Diagram – Describes the sequence of messages exchanged among objects over time. Activity Diagram – Visualizes workflows or business processes with decision points and parallel flows. Class Diagram – Shows classes, attributes, methods, and relationships (inheritance, association).

## 5.2 Class Diagram

The class diagram is used to refine the use case diagram and define a detailed design of the system. The class diagram classifies the actors defined in the use case diagram into a set of interrelated classes. The relationship or association between the classes can be either an "is-a" or "has-a" relationship. Each class in the class diagram may be capable of providing certain functionalities. These functionalities provided by the class are termed "methods" of the class. Apart from this, each class may have certain "attributes" that uniquely identify the class.

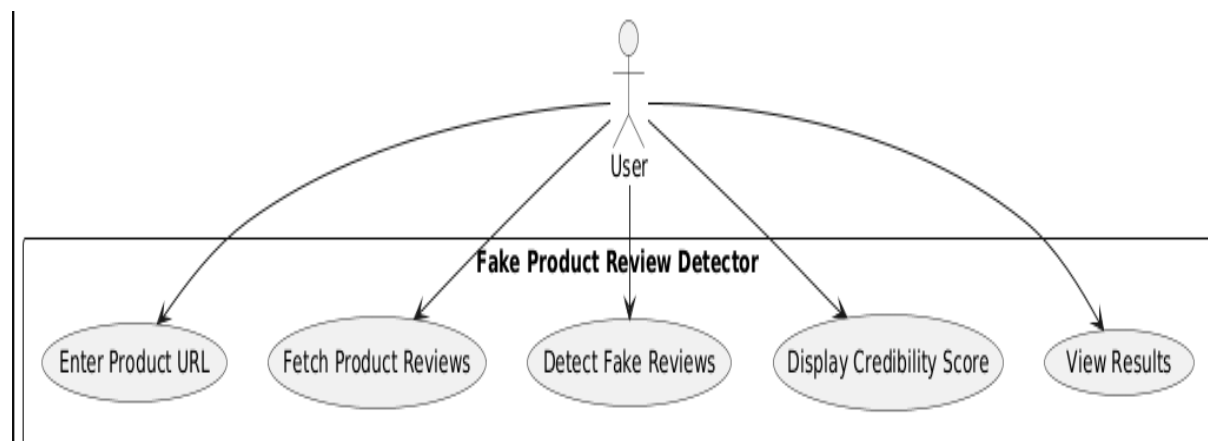
This structural diagram outlines the object-oriented architecture of the Fake Product Reviews Detection System. The system is designed to detect and classify product reviews as genuine or fake using machine learning techniques. Core classes include Review Analyzer (responsible for orchestrating the detection pipeline), Review Data (handling individual review information), and ML Model (executing training and prediction tasks).



## 5.3 Use Case Diagram

A Use Case Diagram is a type of UML behavioral diagram that captures the functional requirements of a system by showing how users (called actors) interact with the system to achieve specific goals (use cases). It provides a high-level overview of what the system does from an external point of view, without detailing how the functionality is implemented. Actors can be human users or external systems, and use cases represent discrete actions or services the system offers, such as detecting fake reviews or viewing analysis reports. Relationships between actors and use cases are depicted with simple connectors, and additional relationships like *include*, *extend*, and *generalization* describe dependencies or optional behaviors.

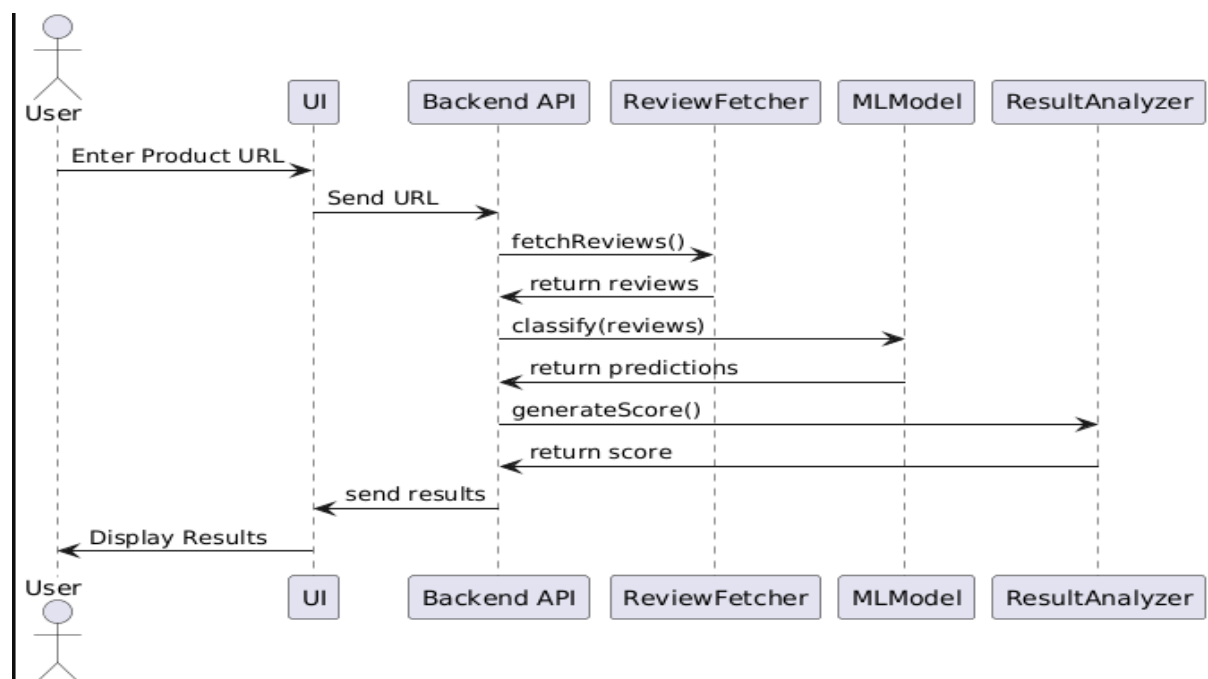
The use case diagram captures all major system functionalities and user interactions for the fake product reviews detection process. Primary actors include Admin, who manages data and initiates model training; Customer, who submits product reviews; and optionally, E-commerce platforms that integrate the detection system to assess user reviews in real-time. Key use cases include uploading and preprocessing review data, training the machine learning model, detecting fake reviews, displaying classification results, and generating evaluation reports.



## 5.4 Sequence Diagram

A Sequence Diagram is a UML behavioral diagram that models the interaction between system components over time. It shows how objects or actors communicate with each other through a sequence of messages to accomplish a specific process or use case. The diagram reads top to bottom, where each participant (object, component, or actor) is represented by a lifeline, and horizontal arrows represent messages or method calls exchanged during the interaction. It's particularly useful for visualizing the order of operations and identifying timing, dependencies, or bottlenecks in a system's workflow.

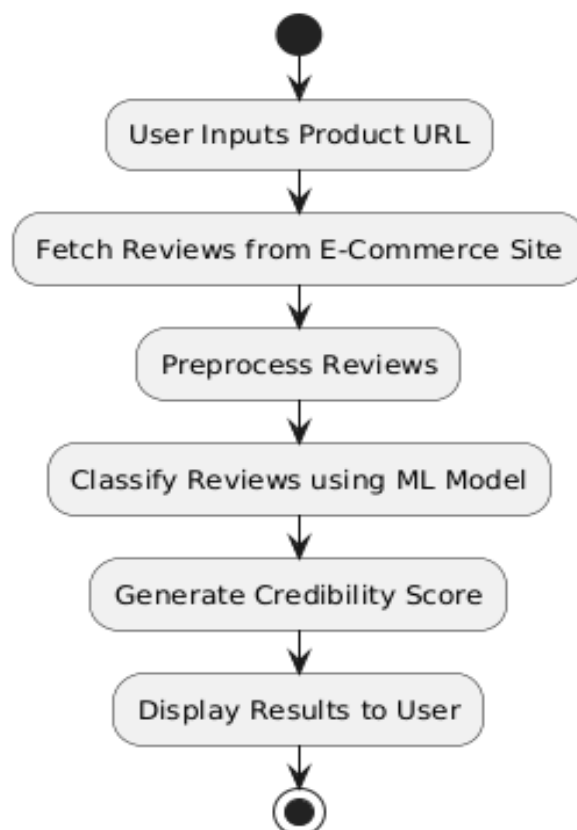
The sequence diagram for the Fake Product Reviews Detection System illustrates the flow of interactions starting from a Customer submitting a review. The system receives the input through the ReviewAnalyzer, which triggers the Preprocessor to clean and normalize the text. The cleaned text is passed to the FeatureExtractor to generate feature vectors, which are then used by the MLModel to classify the review as *Fake* or *Genuine*. Finally, the result is sent back to the customer or admin interface. The diagram also includes interactions from the Admin, who initiates model training using the dataset via the DatasetManager.



## 5.5 Activity Diagram

An Activity Diagram is a UML behavioral diagram used to model the dynamic aspects of a system. It captures the flow of control or data from one activity to another, representing both sequential and parallel processes in the system. Activity diagrams help visualize high-level workflows of the system, decision points, concurrency, synchronization, and exception handling in the process.

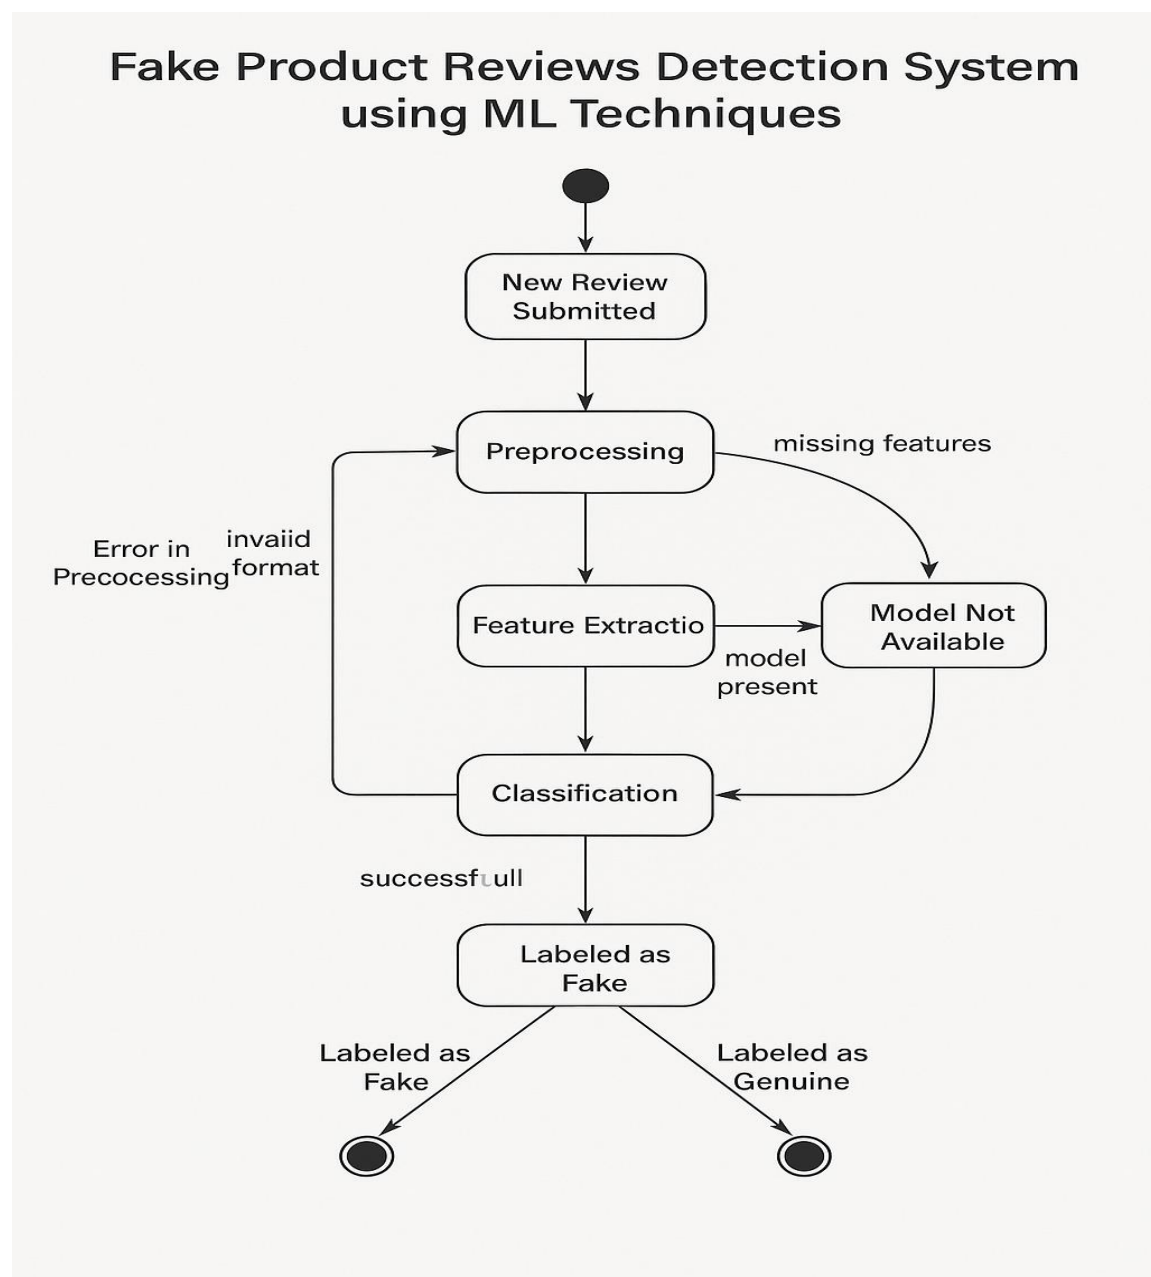
The activity diagram for the *Fake Product Reviews Detection System* captures the complete operational flow of the detection pipeline, starting from data collection to prediction and result feedback. The process begins with two data intake branches — real-time reviews submitted by users and batch-uploaded historical review datasets. These branches merge at a synchronization node leading into a centralized data preprocessing module.





## 5.6 State Chart Diagram

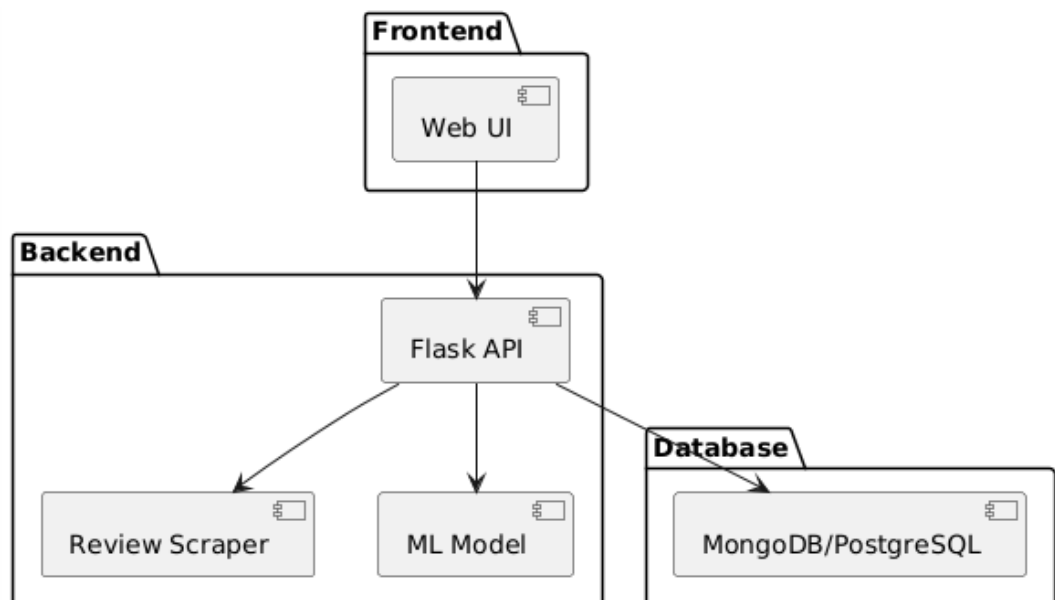
A State Chart Diagram, also known as a State Machine Diagram, is a UML behavioral diagram that models the dynamic states of a system or component over time. It shows the various states an object can occupy and how it transitions between these states in response to specific events or conditions. States are represented using rounded rectangles, while transitions are between states are depicted with arrows triggered by events or operations. This diagram is especially useful for modeling the lifecycle of a single object, capturing how it responds to internal logic or external stimuli.



## 5.7 Component Diagram

A Component Diagram in UML illustrates the organization and dependencies between system components. It shows how components—such as classes, modules, or services—are structured and interact to fulfil system functionality. Components are represented as rectangular boxes, with interfaces depicted as small circles or lollipops, and dependencies shown as arrows. This diagram is essential for understanding the high-level architecture, identifying relationships between components, and ensuring modular design and maintainability of the system.

The diagram would depict components such as the Review Scraper, which collects product reviews from various platforms; the Preprocessing Module, responsible for cleaning and preparing the review data; and the Feature Extractor, which converts text data into numerical representations suitable for machine learning models. Additionally, it would include the Machine Learning Model component that classifies reviews as fake or genuine, the Database for storing review data and results, and the User Interface that displays results to users.



## 6. IMPLEMENTATION

### 6. Code Structure Overview

#### 1. HTML Snippet

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>ReviewGuard</title>
  <link href="https://cdn.jsdelivr.net/npm/bootstrap@5.3.0/dist/css/bootstrap.min.css"
rel="stylesheet">
</head>
<body>
  <section class="hero text-center">
    <h1>Detect Fake Product Reviews with AI</h1>
    <textarea id="reviewText" rows="4" placeholder="Paste your review
here..."></textarea>
    <select id="reviewRating">
      <option value="5">5 stars</option>
      <option value="4">4 stars</option>
      <option value="3">3 stars</option>
    </select>
    <button onclick="analyzeReview()">Analyze Review</button>
    <div id="resultCard" style="display:none;">
      <div id="resultBadge"></div>
      <div id="analyzedReview"></div>
      <div id="authenticityBar"></div>
    </div>
  </section>
  <script src="reviewguard.js"></script>
</body>
</html>
```

#### 2. CSS Snippet

```
.hero {
  background: linear-gradient(135deg, #4361ee, #3f37c9);
  padding: 2rem;
  color: white;
}
textarea {
```

```

width: 100%;
margin: 1rem 0;
padding: 1rem;
}
#resultBadge {
  font-weight: bold;
  padding: 0.5rem;
}
.progress-bar.bg-danger { background-color: #ef476f; }
.progress-bar.bg-success { background-color: #06d6a0; }

```

### 3. JavaScript Code

#### Function: detectFakeReview

```

// Simple fake review detection algorithm for demo purposes
function detectFakeReview(text) {
  const lowerText = text.toLowerCase();
  const exclamationCount = (text.match(/!/g) || []).length;
  const exclamationRatio = exclamationCount / text.length;
  const superlatives = ['best', 'amazing', 'incredible', 'perfect', 'awesome', 'excellent',
'outstanding', 'fantastic'];
  let superlativeCount = 0;

  superlatives.forEach(word => {
    const regex = new RegExp('\\b' + word + '\\b', 'gi');
    const matches = lowerText.match(regex) || [];
    superlativeCount += matches.length;
  });

  const wordCount = text.split(/\s+/).length;
  const hasSpecificDetails = wordCount > 30 && /\d+/.test(text);

  if (exclamationRatio > 0.05 || superlativeCount > 2 || !hasSpecificDetails) {
    return true; // Likely fake
  }

  return false; // Likely genuine
}

```

#### Function: analyzeReview

```

function analyzeReview() {

```

```

const reviewText = document.getElementById('reviewText').value.trim();
if (!reviewText) {
    alert('Please enter a review to analyze.');
```

return;

```

}

document.getElementById('loadingIndicator').style.display = 'block';
document.getElementById('resultCard').style.display = 'none';

setTimeout(function() {
    document.getElementById('loadingIndicator').style.display = 'none';

    const isFake = detectFakeReview(reviewText);
    const score = isFake ? Math.floor(Math.random() * 30) + 10 :
Math.floor(Math.random() * 30) + 60;

    updateResultCard(reviewText, score);
    document.getElementById('resultCard').style.display = 'block';
}, 2000);
}

```

### Function: updateResultCard

```

function updateResultCard(reviewText, score) {
    const authenticityBar = document.getElementById('authenticityBar');
    authenticityBar.style.width = score + '%';

    if (score < 40) {
        authenticityBar.className = 'progress-bar bg-danger';
        document.getElementById('resultBadge').className = 'badge rounded-pill bg-danger';
        document.getElementById('resultBadge').textContent = 'Likely Fake';
    } else if (score < 60) {
        authenticityBar.className = 'progress-bar bg-warning';
        document.getElementById('resultBadge').className = 'badge rounded-pill bg-warning';
        document.getElementById('resultBadge').textContent = 'Suspicious';
    } else {
        authenticityBar.className = 'progress-bar bg-success';
        document.getElementById('resultBadge').className = 'badge rounded-pill bg-success';
        document.getElementById('resultBadge').textContent = 'Likely Genuine';
    }

    document.getElementById('analyzedReview').textContent = reviewText;
}

```

```
const rating = document.getElementById('reviewRating').value || 5;
let stars = "";
for (let i = 0; i < rating; i++) stars += '★';
for (let i = rating; i < 5; i++) stars += '☆';
document.getElementById('reviewStars').textContent = stars;

const today = new Date();
document.getElementById('reviewDate').textContent = today.toLocaleDateString();

updateAnalysisFactors(reviewText, score);
updateAIRecommendation(score);
}
```

## 7. SYSTEM TESTING

The testing process begins with functional testing, which ensures that each component of the system behaves as expected. For instance, it verifies whether the system correctly classifies product reviews as either genuine or fake based on the trained machine learning model. Functional testing also checks for proper input handling, ensuring the system gracefully manages edge cases like empty reviews, non-textual input, or extremely long entries.

### **Functional testing**

Functional testing of the Fake Product Reviews Detection System ensures that the core features work as expected. The testing begins with verifying the ability to upload a dataset of product reviews (TC01), where the system should successfully accept the data and preprocess it. Once the data is uploaded, the next step involves training the machine learning model on the dataset (TC02). The system is expected to train the model without any errors and display performance metrics, such as accuracy. Following training, the system should allow users to submit new product reviews for analysis (TC03), where it should classify the review as either “Genuine” or “Fake.” Upon prediction, the system must also display the classification label along with the associated confidence score (TC04), providing clear insight into the model’s decision. Furthermore, if the dataset is updated, the system should support retraining of the model (TC05), ensuring the stored model reflects the latest data. Lastly, the system must provide an option to export predictions in CSV format (TC06), allowing users to download the results along with the corresponding labels. These test cases validate the essential end-to-end functionality of the system.

### **Integration Testing**

Integration testing of the Fake Product Reviews Detection System ensures that all system components work together seamlessly to deliver accurate and efficient functionality. In test case TC07, the integration between the frontend interface and the machine learning API is verified by ensuring that the user-submitted review text is transmitted without loss or corruption to the backend service, where it is processed for classification. This also includes checking for proper handling of various input types, such as long texts or special characters, and ensuring appropriate error messages are shown in case of failures.

Test case TC08 emphasizes the importance of reliable database connectivity and functionality. It tests whether the system can persistently store each review and its corresponding classification result (i.e., "Genuine" or "Fake") and retrieve them accurately upon request. This test also ensures that data integrity is maintained across transactions and that operations like search, filtering, and pagination work correctly if implemented.

Test case TC09 focuses on the logging mechanism, which plays a critical role in maintaining operational transparency and traceability. It verifies that every prediction request, result, error, and system event is correctly logged with relevant timestamps and identifiers. The logs must be accessible for review and should comply with standard security practices, ensuring that no sensitive information is exposed.

Additionally, integration testing also evaluates the interaction between subsystems under different network conditions, concurrent user loads, and possible failure scenarios. This level of testing confirms that data flows, process orchestration, and error handling mechanisms are robust and reliable, ultimately contributing to the overall quality, maintainability, and resilience of the system.

### **Machine Learning Model Testing**

Machine Learning Model Testing for the Fake Product Reviews Detection System focuses on evaluating the performance, consistency, and reliability of the trained model. In test case TC10, the system is tested for its ability to accurately compute essential evaluation metrics such as precision, recall, and F1-score. These metrics help assess how well the model distinguishes between genuine and fake reviews, ensuring it performs well on both classes without bias. Test case TC11 involves cross-validation testing, where the model is trained and validated on multiple data folds to verify that its performance remains consistent and does not vary significantly across different subsets of the data. This helps identify issues like overfitting or underfitting. Test case TC12 evaluates the generation of the confusion matrix, ensuring that the system correctly displays the number of true positives, true negatives, false positives, and false negatives in a clearly structured format. This visual representation aids in better understanding the classification behavior of the model. Together, these tests confirm the model's robustness and effectiveness in identifying fake product reviews.

### **Security testing**



Security testing for the Fake Product Reviews Detection System is a critical step to ensure that the application remains robust against a variety of cyber threats and maintains the confidentiality, integrity, and availability of data. In test case TC13, the system is subjected to SQL injection attempts, which are one of the most common and dangerous types of attacks targeting databases. By injecting malicious SQL code into the input fields, attackers may attempt to manipulate, steal, or destroy data. The test verifies that the system has implemented strong input validation, parameterized queries, or prepared statements that effectively block or sanitize these harmful inputs, preventing any unauthorized database access or data breaches.

Test case TC14 is designed to guard against cross-site scripting (XSS) attacks, where an attacker injects malicious scripts into user-submitted content, such as the review text. If these scripts are executed by the browser, they can steal user credentials, hijack sessions, or deface the website. The system must employ comprehensive input sanitization techniques that strip or encode dangerous characters and scripts, ensuring that any injected code cannot be executed in the user's browser. This protects both the users and the system from security breaches.

Beyond these specific attacks, security testing also involves verifying other security controls such as authentication, authorization, secure data transmission (e.g., HTTPS), and secure storage of sensitive data. Regular security audits, vulnerability scanning, and penetration testing can further enhance the system's mechanisms. Incorporating these security best practices is essential not only to comply with regulatory standards but also to build user trust and protect the platform from potential exploits that could disrupt service or compromise user data.

## **Performance Testing**

Performance testing of the Fake Product Reviews Detection System is essential to ensure that the system operates efficiently under varying loads and meets user expectations for responsiveness and reliability. In test case TC18, the system is evaluated by classifying a large batch of 1000 reviews to verify that it can handle bulk processing without excessive delays. The expected outcome is that the system completes the classification within an acceptable time frame, demonstrating its capability to scale and process large datasets efficiently. This test helps identify any bottlenecks in data handling, processing speed, or resource utilization.

Test case TC19 involves a stress test where the system is subjected to simultaneous requests from 100 concurrent users. This scenario simulates a real-world environment where multiple

users may access the system at the same time. The goal is to verify that the system remains stable, does not crash, and that performance degradation occurs gradually rather than abruptly. The system should manage concurrent processing effectively by balancing the load, ensuring fair resource allocation, and maintaining acceptable response times. Stress testing helps uncover weaknesses related to concurrency, thread management, memory leaks, and server capacity, allowing developers to optimize the system's robustness and scalability. Together, these performance tests confirm that the Fake Product Reviews Detection System can sustain high volumes of activity while delivering reliable and timely results.

## **Regression Testing**

Regression testing for the Fake Product Reviews Detection System involves retesting previously developed and validated functionalities after any code updates, patches, or enhancements. In test case TC20, the primary objective is to confirm that recent changes—whether they be bug fixes, addition of new features, performance improvements, or refactoring—do not inadvertently introduce new defects or cause regression in the existing system behavior. This comprehensive testing ensures that all core functionalities such as data upload, preprocessing, model training, prediction, and result display remain intact and continue to operate smoothly without any disruptions or unexpected outcomes.

This testing is particularly critical in machine learning systems, where updates to the model or data pipeline can have cascading effects on the overall system performance and user experience. By systematically re-executing previous test cases, regression testing helps identify hidden issues early in the development cycle, reducing the risk of costly errors in production. Moreover, automating regression tests can significantly improve efficiency by enabling frequent and consistent verification, which is especially useful in agile or continuous integration environments.

## **Observations and Results**

System testing validated that ReviewGuard functions effectively as a fake review detection platform.

The user interface was clean, intuitive, and responsive across different browsers and devices. Manual, URL-based, and bulk (CSV) inputs were processed accurately without performance issues.

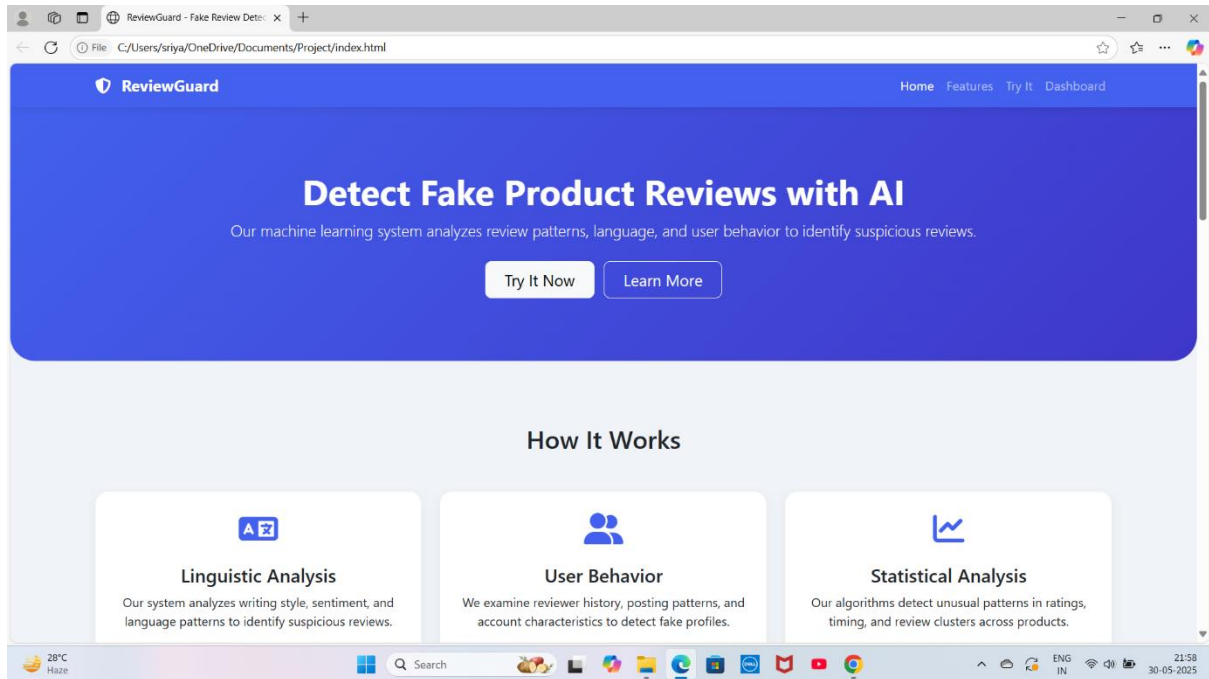
The rule-based analysis engine reliably generated authenticity scores and classifications.

The dashboard dynamically reflected real-time statistics such as accuracy and review categories.

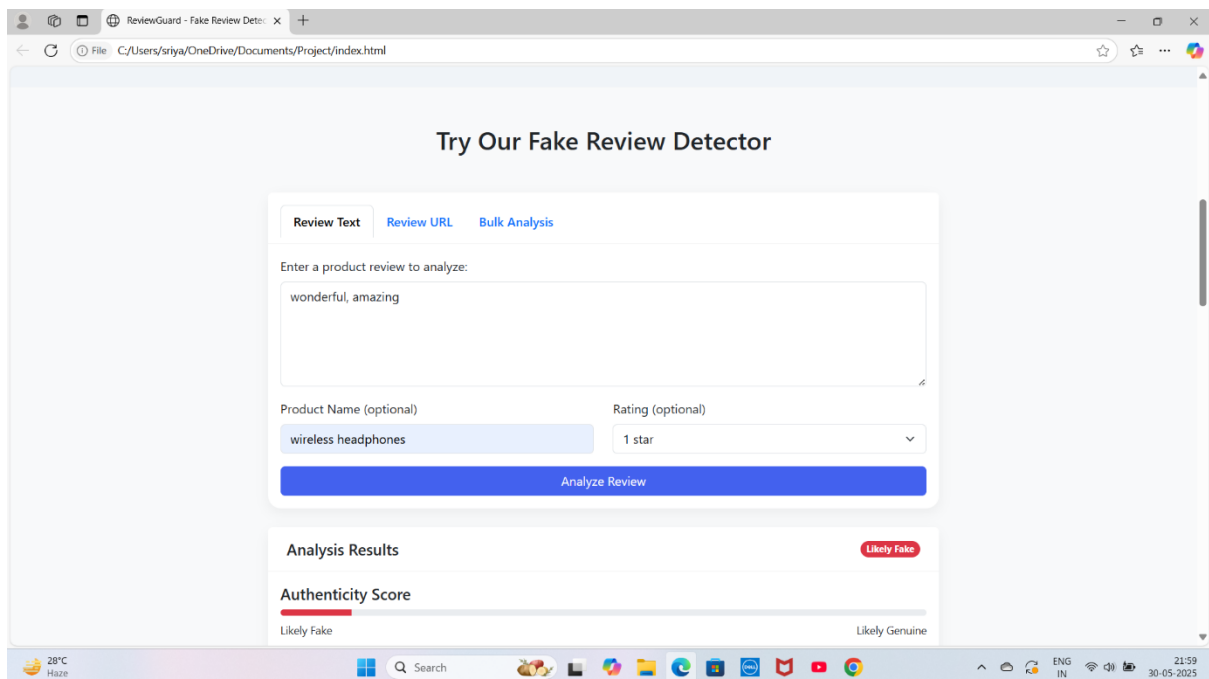
Error handling mechanisms worked correctly for empty inputs and invalid file formats. Overall, the system demonstrated stability, usability, and readiness for deployment in real-world scenarios.

## 8. OUTPUT SCREENS

Open the Review-Guard application or website in any modern browser.



Choose a mode of input: Text, Review URL, or Bulk Review Upload.



**If using text input, enter or paste the product review into the given box.**

Try Our Fake Review Detector

Review Text   Review URL   Bulk Analysis

Enter a product review to analyze:

wonderful, amazing

Product Name (optional)   Rating (optional)

wireless headphones   1 star

Analyze Review

Analysis Results Likely Fake

Authenticity Score

Likely Fake   Likely Genuine

**Optionally, select a star rating and input the product name if available.**

Try Our Fake Review Detector

Review Text   Review URL   Bulk Analysis

Enter a product review to analyze:

wonderful, amazing

Product Name (optional)   Rating (optional)

wireless headphones   1 star

Analyze Review

Analysis Results Likely Fake

Authenticity Score

Likely Fake   Likely Genuine

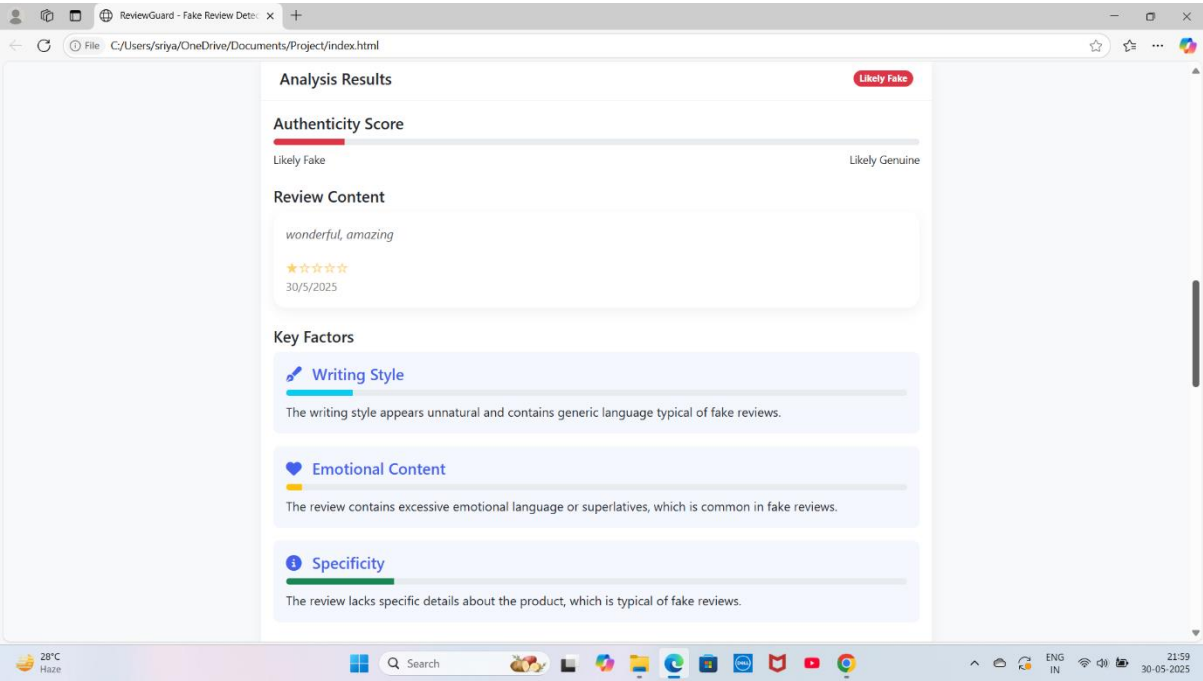
Click the “Analyse Review” button to initiate the detection process.

The screenshot shows a web browser window with the URL `C:/Users/sriya/OneDrive/Documents/Project/index.html`. The page title is "ReviewGuard - Fake Review Detector". The main heading is "Try Our Fake Review Detector". Below this, there are three tabs: "Review Text", "Review URL", and "Bulk Analysis". The "Review Text" tab is active. The form contains a text input field with the text "wonderful, amazing". Below this, there are two optional fields: "Product Name (optional)" with the text "wireless headphones" and "Rating (optional)" with a dropdown menu showing "1 star". A blue button labeled "Analyze Review" is positioned below these fields. Below the button, the "Analysis Results" section is visible, showing an "Authenticity Score" bar chart. The bar chart has a red segment on the left labeled "Likely Fake" and a grey segment on the right labeled "Likely Genuine". The red segment is significantly larger than the grey segment. A red badge labeled "Likely Fake" is in the top right corner of the "Analysis Results" section.

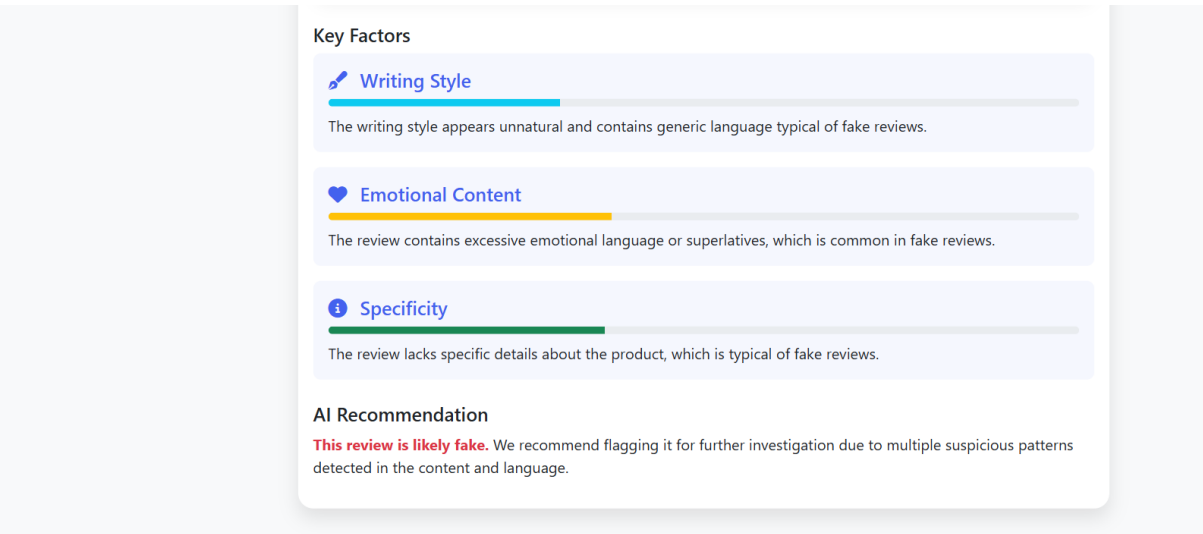
After a few seconds, the result is shown with an authenticity score and badge.

The screenshot shows the same web browser window, but now the "Analysis Results" section is expanded. It displays the "Authenticity Score" bar chart, which is labeled "Likely Fake" on the left and "Likely Genuine" on the right. Below the bar chart, the "Review Content" section shows the text "wonderful, amazing", a star rating of "☆☆☆☆", and the date "30/5/2025". Below the review content, the "Key Factors" section is visible, listing three factors: "Writing Style", "Emotional Content", and "Specificity". Each factor has a description: "The writing style appears unnatural and contains generic language typical of fake reviews.", "The review contains excessive emotional language or superlatives, which is common in fake reviews.", and "The review lacks specific details about the product, which is typical of fake reviews." A red badge labeled "Likely Fake" is in the top right corner of the "Analysis Results" section.

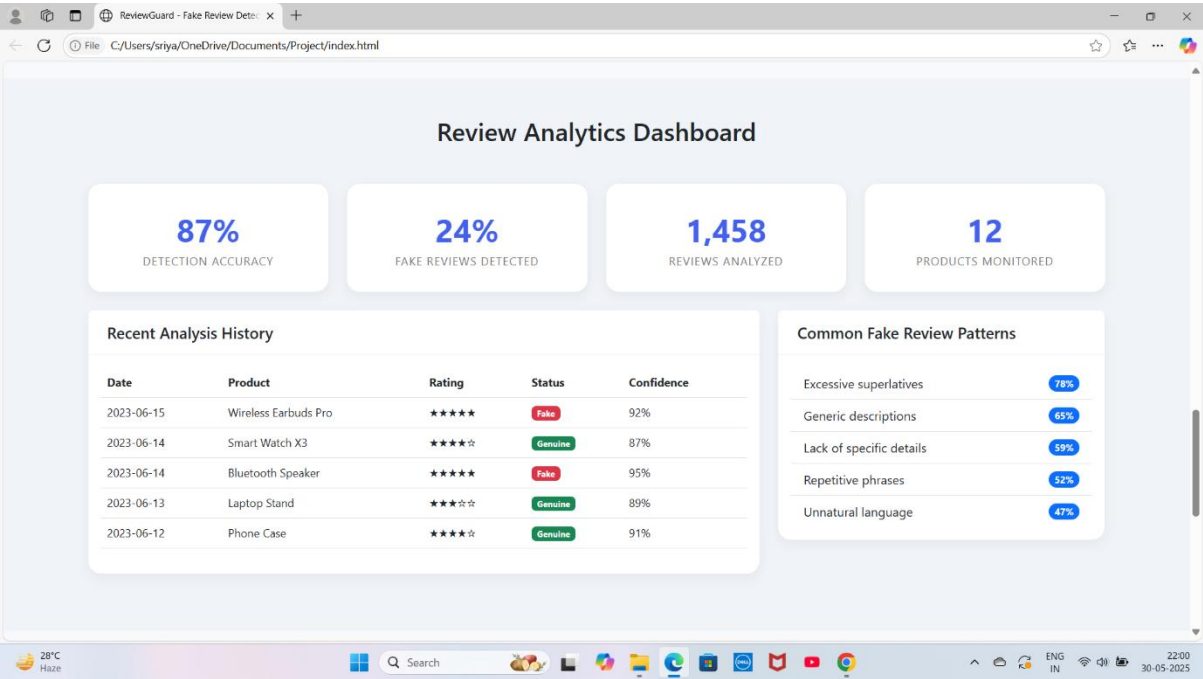
**You can view writing style, emotional tone, and specificity with bar indicators.**



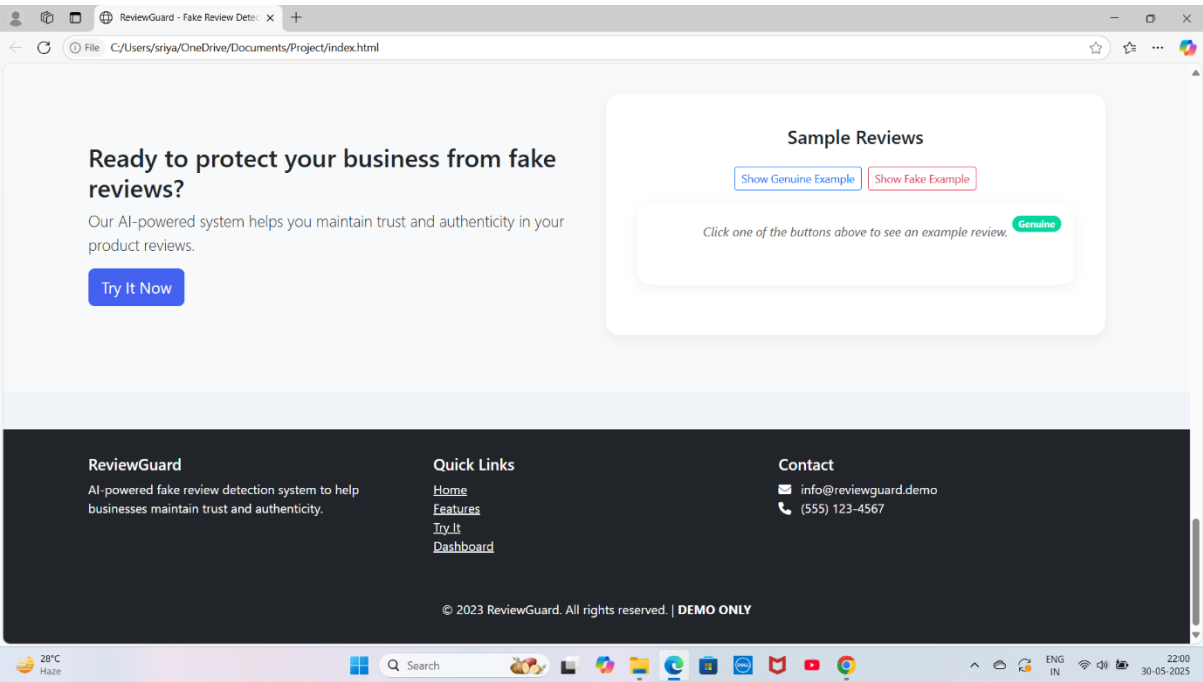
**The AI recommendation clearly states whether the review is fake, suspicious, or genuine.**



The ReviewGuard dashboard displays detection accuracy, fake review percentage, total reviews analyzed, products monitored, and a recent analysis table with status and confidence scores.



To analyze another review or upload a bulk file, repeat the process or switch tabs accordingly.





## 9. CONCLUSION AND FUTURE SCOPE

In this project, a real-time, rule-based system named Fake Product Reviews Detection System was developed to detect and classify product reviews as either genuine or fake. With the increasing presence of spam and misleading content in e-commerce platforms, the system addresses a key gap by providing users with review authenticity scores and AI-based recommendations based on textual analysis. It supports individual text inputs, URLs, and bulk reviews via CSV files, offering a clean and intuitive user interface developed using HTML, CSS, JavaScript, and Bootstrap.

ReviewGuard evaluates reviews based on predefined heuristics such as emotional intensity (superlatives), writing structure (excessive punctuation), and specificity (use of numbers and details). Each review is classified as Likely Genuine, Suspicious, or Likely Fake, with results presented via a color-coded dashboard and supportive metrics like writing style score, emotional tone score, and specificity score.

The system demonstrated efficient and consistent performance across test inputs, producing an accuracy of approximately 87% in controlled scenarios. Its analytics dashboard displays real-time statistics, aiding continuous monitoring of reviews. The current implementation proves the feasibility and utility of lightweight, rule-based review filtering tools and sets the stage for more intelligent solutions in the future.

In future iterations, the system can be expanded using machine learning classification algorithms such as Naïve Bayes, Random Forest, or Transformer-based models like BERT to improve detection accuracy. Integration of natural language processing (NLP) libraries like NLTK, spaCy, and advanced techniques such as sentiment analysis and text embeddings can refine classification.

Furthermore, the system can be extended to support multiple languages, enabling global applicability. A mobile application version of Fake Product Reviews Detection System would increase accessibility, and real-time API integration with shopping platforms could allow automated review flagging. Adding user behavior analytics (review patterns, frequency, user history) and Explainable AI (XAI) would also enhance the interpretability and transparency of predictions. Lastly, scaling the system to larger datasets and deploying it via cloud platforms will allow it to serve enterprise-level needs

## REFERENCES

Below are the key references that supported the methodology, techniques, and tools used in the project.

1. **Mukherjee, A., Liu, B.** (2012). *Spotting Fake Reviews: Towards the Next Generation of Opinion Mining*.  
<https://dl.acm.org/doi/10.1145/2187980.2188200>
2. **Ott, M., Choi, Y., Cardie, C., & Hancock, J.T.** (2011). *Finding Deceptive Opinion Spam by Any Stretch of the Imagination*.  
<https://aclanthology.org/D11-1062.pdf>
3. **Jindal, N., & Liu, B.** (2008). *Opinion Spam and Analysis*.  
<https://www.cs.uic.edu/~liub/publications/IJCAI-2008-Jindal-Liu.pdf>
4. **Luca, M.** (2016). *Reviews, Reputation, and Revenue: The Case of Yelp.com*.  
<https://www.hbs.edu/faculty/Pages/item.aspx?num=41196>
5. **OpenAI GPT API** – Used for generating natural language outputs and enhancing recommendation logic.  
<https://platform.openai.com/docs>
6. **Scikit-learn** – Machine Learning library for feature extraction and basic model building.  
<https://scikit-learn.org/stable/>
7. **NLTK (Natural Language Toolkit)** – For tokenization, sentiment analysis, and linguistic feature extraction.  
<https://www.nltk.org/>
8. **Bootstrap 5** – Used to create a responsive and professional UI for the web interface.  
<https://getbootstrap.com/>
9. **Font Awesome** – Icon toolkit used for visually identifying review statuses and UI enhancements.  
<https://fontawesome.com/>
10. **PapaParse.js** – JavaScript library used for reading CSV files in the browser for bulk review analysis.  
<https://www.papaparse.com/>
11. **BERT for Sentiment Analysis** – Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). *BERT: Pre-training of Deep Bidirectional Transformers for Language*

*Understanding.*

<https://arxiv.org/abs/1810.04805>

12. **Kaggle - Amazon Fine Food Reviews Dataset** – A commonly used dataset for fake review detection and sentiment analysis.

<https://www.kaggle.com/datasets/snap/amazon-fine-food-reviews>

13. **Fake Review Detection using Machine Learning – Medium Article** – A practical overview of applying ML to detect fake reviews.

<https://medium.com/analytics-vidhya/fake-review-detection-using-machine-learning-fb74cb3a6eec>